


Image Encoder LoRA — Verified Final Configuration (No Local Hyperparameter Sweep Recorded)

Parameter-efficient vision adaptation

Frozen CLIP ViT-B/32 weights

Trainable LoRA adapters · rank=32 · alpha=32

**All-linear targets · q/k/v/out · fc1/fc2
· visual projection**



LoRA dropout=0.0 · cosine scheduler · weight decay 0.05 · 25 epochs · fixed final checkpoint